

The Fog of Innovation

Finding Clarity at the Intersection of Discovery and Deployment

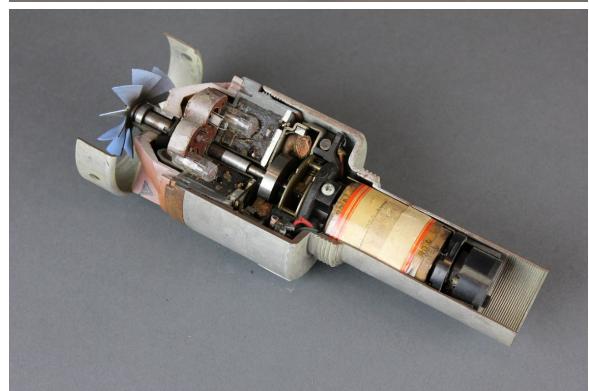
TYLER D. DOLEŽAL¹

World War II is often remembered for its battles, advancements in military tactics, and deep loss, and yet it also marked a turning point in the relationship between science and warfare. Radar, the atomic bomb, synthetic rubber—these weren't just technological breakthroughs; they were decisive forces that shaped the outcome of the war and redefined what science could mean in a time of national urgency. During that era, scientific discovery and military application were deeply intertwined, with end-users and researchers working in lockstep under urgent timelines and shared stakes.

Fast forward to today. We live in another era of global uncertainty—an era that once again demands innovation. But something feels different. As an acquisition officer in the United States Air Force and a researcher in materials science, I stand at the intersection of discovery and defense. On one hand, I see a dazzling spectrum of innovation: high-entropy alloys, self-healing coatings, AI-optimized materials, and nanostructured thermal barriers. On the other, I face the growing challenge of keeping up—of reviewing this surge of work with the depth and discernment it deserves. Each month, thousands of papers are published in my field alone. Many tout revolutionary findings, but few offer the clarity, maturity, or reproducibility needed for actual deployment. And even fewer are written with the end-user in mind—the engineer trying to increase engine life, the maintainer seeking corrosion resistance in harsh climates, or the acquisition officer evaluating which tech is worth a billion-dollar investment.

The problem isn't a lack of brilliance. It's the pace, the noise, and the widening gulf between academic progress and applied utility. We've built an engine of discovery with no brakes and no navigation system. What does that look like in practice? Too often, papers emphasize novelty over reproducibility. They present promising results with no benchmark against existing materials. They lack long-term validation or environmental robustness. And more critically, they rarely speak the language of application—of constraints, cost, or integration into real-world systems.

In the decades following World War II, scientific publishing followed a deliberate rhythm.



(Top) A quartz crystal clock used to coordinate military operations with unprecedented precision.

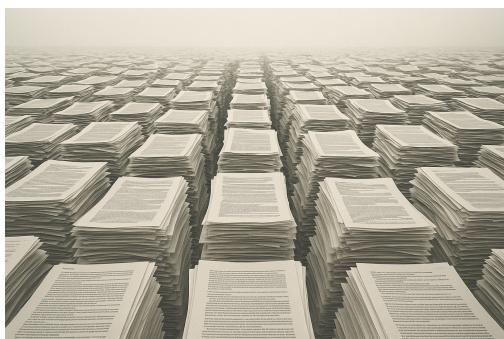
(Bottom) Cutaway of a radio proximity fuse. This early radar-powered sensor triggered detonation near a target.

Photo Source: National Institute of Standards and Technology Digital Archives

¹With ongoing intellectual and editorial support from Nick Valverde, Ph.D., whose insight and soundboarding helped shape this piece.

Research was slow, peer review was methodical, and breakthroughs were rare—but often deeply validated. That model, however, is long gone. Today, we live in the age of more. More journals. More papers. More citations. The accessibility of open-access publishing, the proliferation of niche journals, and the relentless pressure on academics to “publish or perish” have created a research ecosystem optimized for output. But there’s a cost.

For an end-user like myself—someone tasked with identifying, evaluating, and implementing scientific advances for national defense—the sheer volume of literature has become a liability. It’s no longer just about staying informed. It’s about surviving the deluge and trying to find meaningful patterns in a sea of noise. This acceleration in publishing speed may benefit academic careers and journal metrics, but it leaves end-users behind. We are overwhelmed not by a lack of innovation, but by an inability to discern which innovations are real, mature, and deployable.



The volume of published research has surged in recent decades. For end-users, the challenge is no longer access, it's assessment.
Illustration generated with DALL·E, April 2025.

As researchers, program managers, and technical advisors, we’ve all been there—we open a paper that promises a significant breakthrough, one of many that we stumble upon or are linked to each week. The results look compelling, the figures are sharp, and the language radiates confidence. But by the end, we’re left wondering: was this benchmarked against existing solutions? Can it survive outside a lab? Is it scalable, or even implementable? That moment of quiet frustration is familiar to many of us—the realization that while the science may be sound, even exciting, it stops short of being truly useful. The critical details are often missing, obscured, or presented in ways that don’t align with how real-world decisions get made. The result is research that inspires curiosity, but not confidence. It

can be exciting science—but it’s not actionable science.

This is the core of the problem: the signal of practical innovation is getting drowned out by the noise of academic momentum. Many papers are optimized for publication, not translation. Their metrics are chosen for novelty, not relevance. Their conclusions are framed around statistical significance, not real-world performance. It’s my job to combine technical expertise with operational insight—to judge whether a material can survive the salt-laden air of a coastal airbase or endure the thermal cycling of a hypersonic vehicle. In that context, novelty isn’t enough. I don’t need the most exciting coating—I need the most deployable one.

Worse still, the way science is communicated often makes it difficult to even ask the right questions. Crucial details are hidden in supplemental sections or omitted entirely. Comparisons to industry-standard materials are rare. Long-term data on fatigue, oxidation, or environmental durability—if collected at all—are buried beneath easier-to-obtain metrics that say little about real-world performance. Yet I’m expected to make decisions based on this. I have to take these papers, incomplete as they often are, and translate them into justifications—recommendations for investment, for adoption, for risk. In a system that thrives on citation counts and novelty, I’m left trying to make operational calls with academic fragments. In defense, that gap isn’t academic—it’s operational. It’s mission-relevant. We aren’t short on innovation. We’re short on clarity.

There’s a well-known phrase in technology development: the “valley of death.” It refers to the chasm between discovery and deployment—the place where promising ideas go to quietly disappear. Nowhere is this more evident than in the world of materials science.

A new coating might show extraordinary performance in lab-scale tests. A simulation

might predict a novel alloy with remarkable thermal stability or resistance to oxidation at high temperatures. But getting from that point to something you can qualify, produce, field, and sustain? That's where most technologies stall. They're not flawed—they're just unfinished. Not enough data. Not enough scale. Not enough understanding of how the material behaves in real environments or interacts with existing systems. Too many caveats.

The incentives of the research ecosystem don't help. Academic publishing rewards novelty, not longevity. Grant agencies favor emerging topics, not mature ones. And industrial partners often wait until the risk is lower—until someone else has de-risked the development. The result is a landscape filled with breakthroughs that are technically valid, but practically stranded. In the defense community, we feel this keenly. The material that survives high-temperature creep in a turbine blade doesn't just need to perform well—it needs to be manufacturable, inspectable, repairable, and certifiable. It needs to work not just once, but for thousands of cycles, in hostile environments, over decades. We don't just need ideas, we need endurance. And too often, we're left trying to cross that valley with a handful of figures and a hope that someone, somewhere, ran the long-term test.

What we need from scientific research isn't just innovation—it's usability. As end-users, we're not asking for miracles. We're asking for relevance. That means giving us context. Show us how the new material performs against standard baselines. Compare it to what's currently fielded. Let us see how it handles the ugly, real-world conditions that can't be avoided in service: thermal shock, oxidation, fatigue, humidity, corrosion, mechanical wear, and manufacturing variability. We need to know if the synthesis method can scale beyond a lab bench, if the testing conditions reflect realistic service environments, and if the performance claims hold over meaningful timeframes. A breakthrough that lasts 10 hours is different from one that lasts 10,000 cycles. And above all, we need clarity. Not just error bars, but assumptions. Not just success stories, but failure modes. Research that's honest about its limitations is far more valuable than research that hides them. Because we don't operate in ideal conditions—we operate in contested, constrained, and unpredictable ones. In defense, a material isn't viable just because it works—it's viable because it can be manufactured at scale, inspected for integrity, certified for flight, and repaired in the field. Without that, it's just a curiosity. And we can't build capability on curiosity. If we want to close the gap between scientific discovery and real-world deployment, we need to shift how research is framed, shared, and evaluated—especially when the intended end-users operate in defense, aerospace, energy, or other high-consequence domains.

First, we need to incentivize **comparative testing**. Novelty is important, but so is benchmarking. If a new material is truly better, it should be tested against the current standard—not just a blank slate. That comparison, more than any one figure of merit, helps end-users make informed decisions. For example, the Materials Genome Initiative (MGI) promotes standardized property reporting and data sharing to accelerate development and enable direct comparisons across materials. As a key partner, NIST has worked to establish reference datasets, interoperability protocols, and validation tools that support more rigorous benchmarking—laying the groundwork for faster, more reliable translation from discovery to deployment.

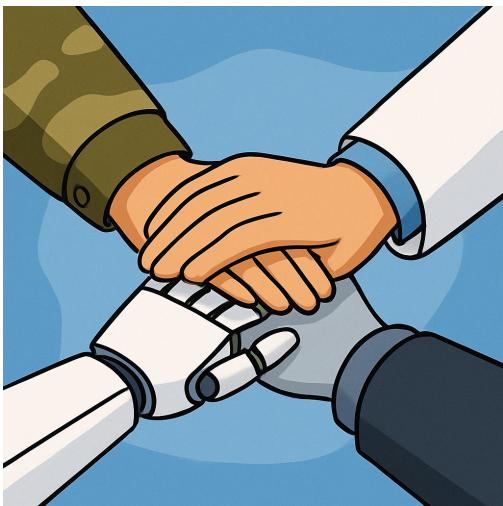


“Just 482 papers left to read before lunch.”

Even in uniform, parsing the scientific firehose under a ticking clock is no small feat for end-users on the operational edge.

Illustration generated with DALL·E, April 2025.

Second, research should include at least some form of **operational framing**. That doesn't mean every paper must simulate a battlefield or turbine blade, but authors can, and should, communicate how their material might behave under realistic conditions: temperature ranges, environmental stressors, manufacturability, and cost. Even a paragraph in the discussion section would go a long way. Agencies like NASA have long adopted this mindset; their Technical Reports Server (NTRS) routinely includes use-case context alongside experimental results.



It takes a coalition: the end-user, the scientist, artificial intelligence, and industry stack hands in a shared mission to turn discovery into deployment.

Illustration generated with DALL·E, April 2025.

Third, we need better tools for **filtering the firehose**. With thousands of papers published each month, it's not about reducing output—it's about helping decision-makers surface what's relevant. Platforms like Semantic Scholar and Connected Papers use AI and graph-based approaches to help researchers and practitioners visualize and prioritize work that's aligned with their needs. Emerging tools are also leveraging large language models (LLMs) to enhance literature searches. For instance, AI2's Paper Finder employs LLMs to mimic human-like reasoning in identifying relevant papers, while Elicit assists in automating research tasks like summarizing papers and extracting data. Additionally, Consensus uses AI to find insights in research papers, aiming to streamline the discovery process for researchers. But while these tools can support better access to operationally relevant findings, they also risk accelerating the churn of research rather than improving its clarity. If we're not careful, the very algorithms meant to help us prioritize impact could instead amplify novelty—further widening the gap between publication and application.

Finally, we must foster more **co-development between researchers and end-users**. Programs that embed operational stakeholders into scientific projects—from conception through validation—are more likely to yield technologies that are both innovative and implementable. The NSF Convergence Accelerator and the long-standing practices of DARPA program managers offer strong models of this embedded collaboration, ensuring science isn't developed in isolation. When researchers understand real constraints, and end-users understand the frontier of what's possible, the result isn't compromise—it's convergence. Science doesn't need to slow down. But it does need to become more legible. More honest. More connected to the missions it's meant to serve.

During World War II, the path from laboratory to field deployment was short, direct, and mission-driven. There wasn't time for ambiguity. Scientists, engineers, and operators worked in sync, because the stakes were clear and the feedback loop was tight. We don't need to recreate the pace of that era—but we do need to recover its sense of purpose. Today, we live in an age of extraordinary scientific capability. But we also live in an age of confusion. A firehose of breakthroughs. A flood of novelty. And for those of us tasked with applying that science to real-world missions—those of us on the sharp end of implementation—the hardest part isn't understanding the technology. It's finding the truth buried underneath the noise. We owe it to ourselves—and to the communities we serve—to reconnect science with its consequences. To prioritize clarity over cleverness. To reward relevance as much as revelation. Because innovation, when it can't be applied, is just potential. And in defense, potential that can't be trusted is a risk we can't afford.

About the Author

Tyler D. Doležal is an active-duty Captain in the United States Air Force and a Ph.D. student in the Department of Materials Science and Engineering at the Massachusetts Institute of Technology. His research focuses on light interstitial doping in nickel-based superalloys. He works at the intersection of fundamental materials science and national defense, with a focus on how emerging technologies transition from lab to field.

Correspondence: dolez120@mit.edu